BPTT step-wise reconstruction: encoder × decoder attention forms

pixel L1 (0-255) · K=64 (196→64 tokens) · 224×224 · ImageNet-100 val (n=5,000) · seed 0

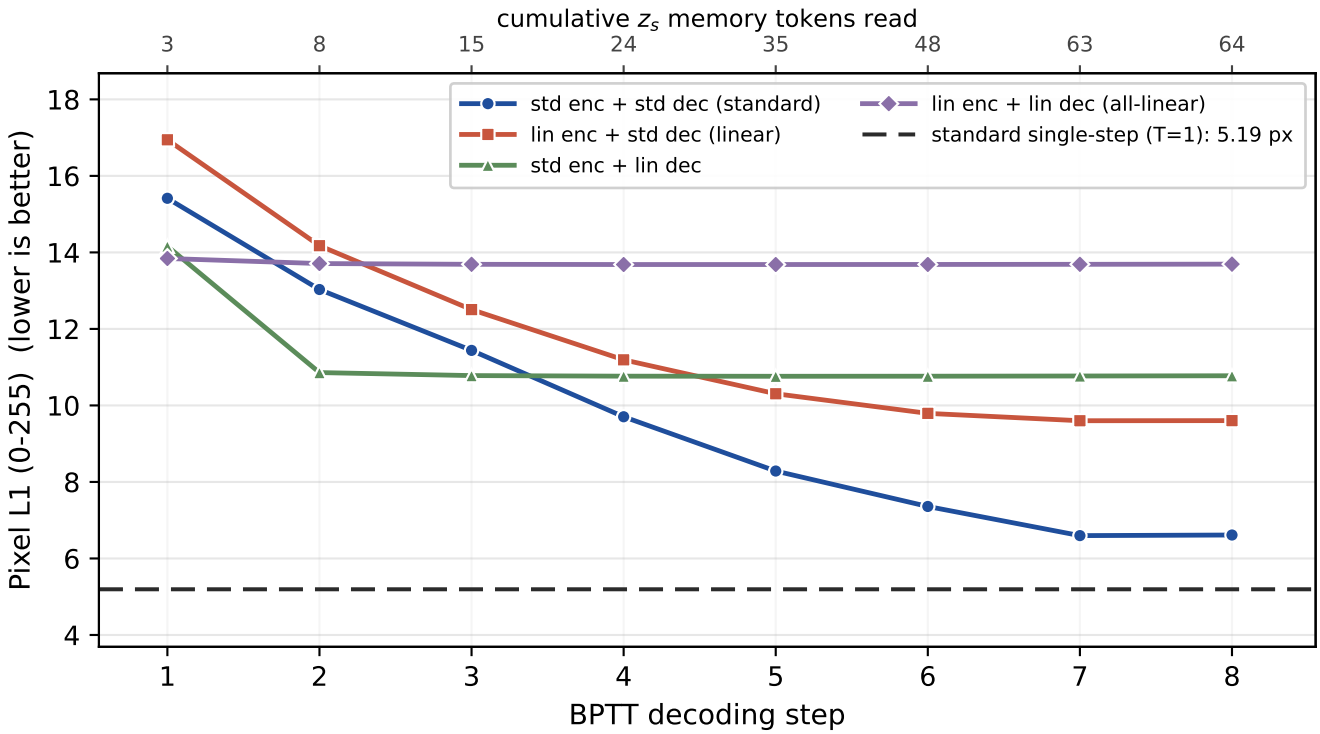

Caveat — training budgets differ: BPTT arms are 6-epoch warm starts (11,406 steps) from their T=1 checkpoints; the T=1 baselines are 45-epoch runs from scratch (42,750 steps). The dashed line is a reference point, NOT a paired ablation. Single seed.