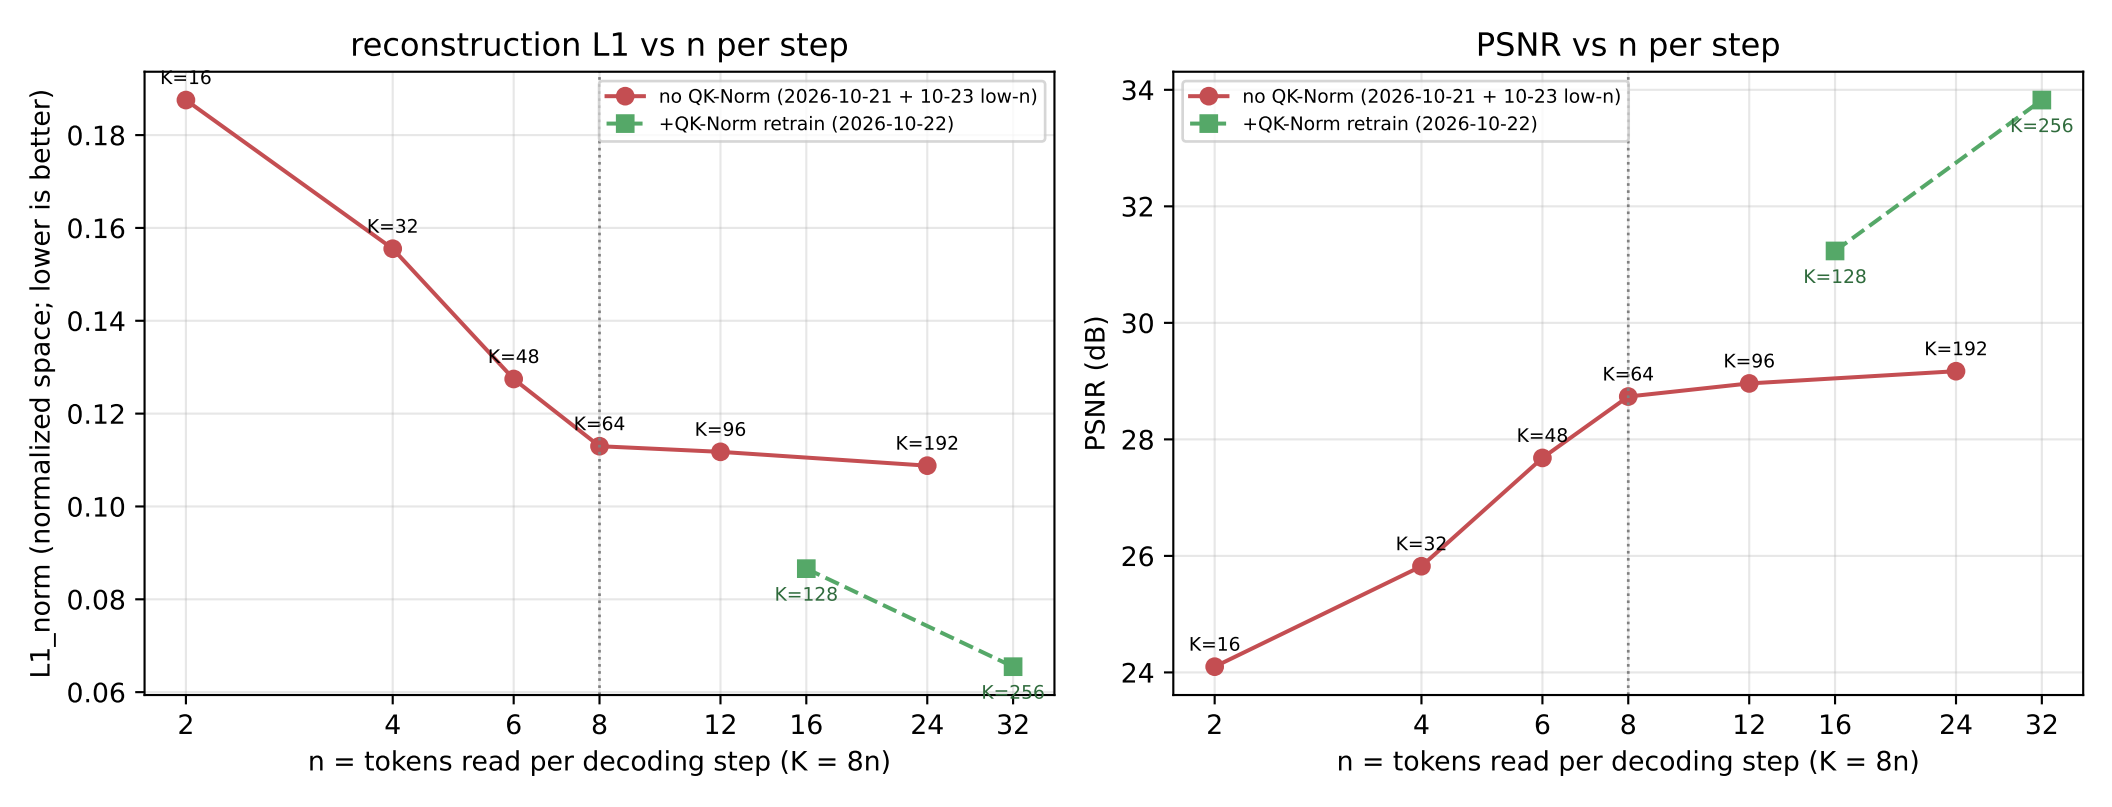
reconstruction L1 vs n per step
L1_norm (normalized space; lower is better)
n = tokens read per decoding step (K = 8n)
no QK-Norm (2026-10-21 + 10-23 low-n)
+QK-Norm retrain (2026-10-22)
K=16
K=32
K=48
K=64
K=96
K=192
K=128
K=256
PSNR vs n per step
PSNR (dB)
n = tokens read per decoding step (K = 8n)