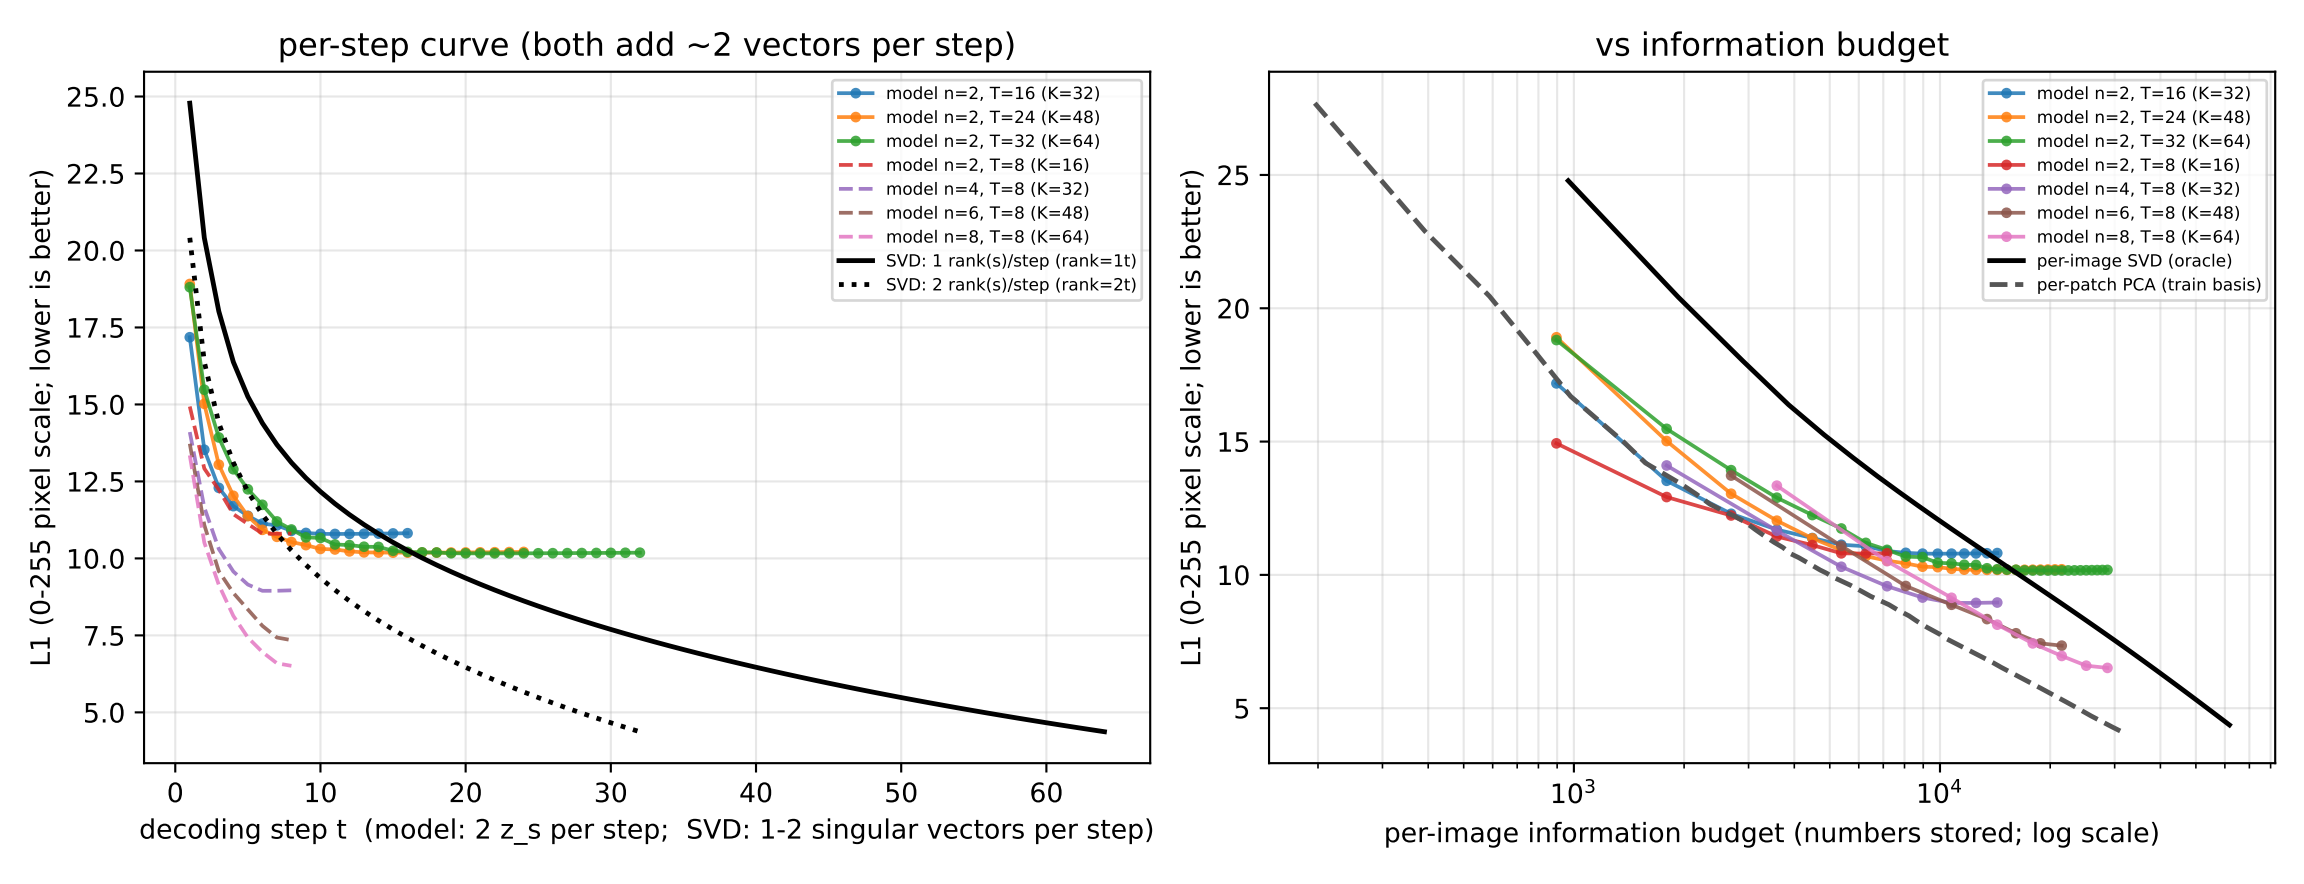
per-step curve (both add ~2 vectors per step)
L1 (0-255 pixel scale; lower is better)
decoding step t (model: 2 z_s per step; SVD: 1-2 singular vectors per step)
model n=2, T=16 (K=32)
model n=2, T=24 (K=48)
model n=2, T=32 (K=64)
model n=2, T=8 (K=16)
model n=4, T=8 (K=32)
model n=6, T=8 (K=48)
model n=8, T=8 (K=64)
SVD: 1 rank(s)/step (rank=1t)
SVD: 2 rank(s)/step (rank=2t)
vs information budget
L1 (0-255 pixel scale; lower is better)
per-image information budget (numbers stored; log scale)
model n=2, T=16 (K=32)
model n=2, T=24 (K=48)
model n=2, T=32 (K=64)
model n=2, T=8 (K=16)
model n=4, T=8 (K=32)
model n=6, T=8 (K=48)
model n=8, T=8 (K=64)
per-image SVD (oracle)
per-patch PCA (train basis)